

Vericoding: Verification-Driven Development from Behavioural Features to Machine-Checked Proofs

The axiomander Toolchain

Gavin Mendel-Gleason
Scidonia

July 2026

Abstract

LLM-driven code generation (“vibecoding”) trades correctness for velocity: generated code may run, but it leaves no durable specification, no behavioural guarantee, and no machine-checkable evidence that it satisfies its intended purpose. The thesis of this work is that **executable behavioural contracts form a semantic intermediate language from which testing, runtime checking, differential validation, and machine-checked proof obligations can all be derived**. We describe a *vericoding* workflow in which a human-authored specification is the primary lasting artifact and generated code is a replaceable implementation detail. We present a toolchain, *axiomander*, to operationalise the approach.

The formal backend targets SNAKELETExN, a small imperative language in the HeapLang tradition extended with first-class exceptions — crucial for modelling Python’s exception-driven control flow. Specifications are stateful: pre/post-conditions, exception arms, and frames are lowered into a Rocq weakest-precondition calculus over a two-cell heap model. The specification comprises Gherkin feature tables, a mathematical contract (pure, boolean-valued predicates over pre/post-state, frames, and exceptions), and an abstract state schema. From this single specification the framework derives executable scenario checks, telemetry assertions, and machine-checked proof obligations in Rocq (Coq) against a hand-verified weakest-precondition calculus. A theory-aware Python-to-SnakeletExn lowerer (§??) produces the service implementation as a SnakeletExn program whose WP refinement is proved against the same contract. A three-valued outcome model (§??) supports constructive counterexample discovery: when an obligation is DISPROVED, the pipeline extracts a concrete runtime witness and certifies a bundling theorem proving the negation.

1 Introduction

Large language models now generate a meaningful fraction of production code. The resulting programs execute without crashing, but leave nothing behind: no specification of intended behaviour, no evidence of correctness, no reproducible conformance regime for the next iteration of the generated code. We call this pattern *vibecoding*: the developer prompts, the model emits, the result runs — and the only long-term artifact is the code itself, with no human-readable statement of what the code is supposed to do.

Vericoding inverts this relationship. In a vericoding workflow, the *specification* is the primary artifact and the code is a replaceable implementation detail. The specification is human-written and human-readable: Gherkin scenarios that domain experts can review, mathematical contracts that state what must hold before and after each operation, and framed state declarations that bound exactly what each operation may read and write. From this single specification every downstream verification activity is derived.




Figure 1: The vericoding pipeline. Features and contracts are the primary lasting artifacts; generated code is a replaceable implementation detail. Every Examples row is exercised against the contract; the same contract lowers to Rocq obligations proved against a hand-verified kernel.

axiomander is a toolchain for vericoding. Its core contract language and theory machinery were extracted from a production citation-checking platform (paperchecker); the invitations example in ?? is taken directly from that codebase.

The toolchain supports two development flows. In *contract-first* development, features are written first, then contracts, then an implementation is built against the contract as a specification, then the implementation is tested and lowered to proof. In *retrospective* development, a developer begins with existing feature files and an existing implementation, writes the contract to capture the behaviour already observed in testing, re-runs the test suite under contract checking to validate the specification, and then proceeds to formal proof. Both flows converge on the same long-lived artifact — the contract — and the same endpoint: a machine-checked proof that the specification is internally sound.

The contract language is *fluid* in the Meyer sense: it is a declarative value assembled from named clauses, not imperative assertions scattered throughout the code. The same contract is both an *executable* test criterion (evaluated as Python booleans against concrete state) and the *input* to machine-checked proof generation (lowered to Rocq propositions). Contracts can be introduced incrementally: testing-only at first, then frames, then exceptions, then invariants, then proof obligations — the toolchain accepts whatever level of specification the domain author provides.

?? summarises the pipeline and ?? walks through a single example step by step.

2 Scientific Contributions

The central thesis of this paper — that a single declarative contract can drive testing, runtime checking, and machine-checked proof — is a realisation of the classical design-by-contract thesis [?, ?], not a new idea in itself. Spec#, Dafny, and JML all pursue variants of this thesis. The contribution of this paper is to show that changes in AI infrastructure now facilitate using this practically *in the large* due to changes in the cost structure of proof. In pursuit of this we also connect up practical mechanisms for connecting the concept of mocks to theories in a way that fuses testing and formal modelling for instance for Python database services.

2.1 Differentially validated library theories

The most distinctive mechanism is the *theory*: a stub for a real library (SQLAlchemy, Python logging, OpenTelemetry) that is given an explicit operational semantics over a pure state model, and that is *differentially validated* against the real library. The same programs are executed against both the stub and the real library, and their observable outputs must agree on the covered fragment; anything outside the fragment is rejected loudly rather than silently approximated. Differential testing is an old technique [?], and adjacent ideas exist in many forms:

CompCert’s external models for I/O [?], QuickChick’s executable semantics for property-based testing [?], K framework’s reference interpreters [?], and ordinary mocking frameworks. **We are unaware of previous work combining executable operational semantics for production libraries with differential validation against the concrete implementation.** It replaces the pervasive but unprincipled use of mocks with a conformance regime that is itself checkable.

2.2 Automatic lowering from Python contracts to Rocq

The lowering pipeline operates in two directions. **Contract lowering** translates Python predicate clauses into SnakeletExn specifications (FunSpecS pre/post over a two-cell heap model); all 23 obligations across four contracts are machine-proved. **Implementation lowering** (§??) translates the Python service function — the SQLAlchemy body inside `reserve()` — into a SnakeletExn program via theory-aware call resolution, then proves WP refinement against the contract’s specification. Why3 extracts to Coq, but from a dedicated specification language [?]; Spec# uses Boogie [?]; Frama-C has its own WP but not Coq back-ends for Python. Our pipeline takes contracts written in a restricted fragment of Python and lowers them into a small imperative language with exceptions (SNAKELETExN — HeapLang extended with first-class exceptions, following [?]), emitting FunSpecS-style obligation files with an automatically proven totality lemma, invariant-preservation lemma, and frame-soundness lemma. The dependency-layered emission for parallel proof generation (§??) is likewise novel.

2.3 Three-valued outcome model

The lowering pipeline does not stop at proving obligations closed. It supports a *three-valued* outcome per obligation: PROVED (the proof closes), DISPROVED (a bundling theorem certifies the negation from a concrete runtime counterexample), and UNKNOWN (neither proof nor disproof is certified). This is enabled by an evidential negation layer (`Lneg.v`), emitted alongside the positive layers when counterexample witnesses are provided, and by per-obligation compile checking that prevents a failed obligation from masking a sound one. The dialectic loop — prove, discover counterexample, certify disproof, fix contract — is the operationalisation of a verification workflow that does not assume the specification is correct a priori. We are unaware of previous automatic lowering pipelines that produce both positive and negative machine-checked verdicts from the same specification.

Foundations in Iris. Our WP calculus is not built from scratch. It is a direct application of the Iris project’s `gen_heap` library [?] and follows the weakest-precondition calculus framework of Iris [?]. The two-cell heap model (`store_loc`, `trace_loc`) is a thin layer of FunSpecS-specific specification logic over Iris’s standard ghost state and points-to reasoning. All propositional-level proof machinery — lookup/insert lemmas, update application, frame soundness — is standard Iris reasoning, reused without modification. Our contribution is the *lowering*: the mechanical translation from Python contracts into Iris-style specification predicates, not the calculus itself.

2.4 Semantic frame discipline with insertion

Frame conditions are old. Eiffel has explicit modifies clauses [?]; Frama-C has `assigns`; Dafny has `modifies`. Our frame system is an engineering refinement: declarative write paths (keyed-row, keyed-attribute, event-channel) generate the complement automatically, eliminating the “everything else unchanged” boilerplate that dominates classical DbC specifications. The keyed-row *insert* discipline — keyed-row write paths may add the args-resolved key, while deletion is always rejected — is a specific rule we have not seen elsewhere. Event-channel frames, covering structured telemetry as first-class frame-citizens, are likewise distinctive.

2.5 Positioning of other machinery

The symmetric projection — one observation function applied before and after execution — is standard practice in program verification [?], not a contribution; we use it because it is the correct discipline. The declarative `SqlDomain` specification that collapses per-domain infrastructure is an engineering ergonomic, not a scientific contribution. The contract language itself — pure, boolean-valued Python predicates — is a design choice motivated by Python familiarity, not a formal advance.

3 Formal Model

Before describing the implementation we fix the semantic objects the toolchain reasons about. The state of a system is partitioned into three strata.

Definition 1 (Specification state). $\Sigma = \langle \mathcal{O}, \mathcal{D}, \mathcal{G} \rangle$ where $\mathcal{O}$ is the observed state (a tuple of typed maps and event-channel tuples read from concrete execution), $\mathcal{D}$ is the derived state (a tuple of pure computations over $\mathcal{O}$, equipped with a consistency check $\text{derive} : \mathcal{O} \rightarrow \mathcal{D}$), and $\mathcal{G}$ is ghost state (auxiliary specification-only state invisible to the implementation).

Definition 2 (Contract satisfaction). A contract $C = \langle \Sigma, \text{pre}, \text{post}, \mathcal{X}, \mathcal{I}, \mathcal{W}, \mathcal{R} \rangle$ is satisfied by an execution $\sigma \xrightarrow{a} \sigma'$ with result r (where a are the call arguments) iff:

- (i) $\text{pre}(\sigma, a)$ holds (admissibility);
- (ii) if no exception condition in $\mathcal{X}$ holds, then $\text{post}(\sigma, a, r, \sigma')$ holds, and every state path outside $\mathcal{W}$ is unchanged between σ and σ' ;
- (iii) if an exception condition $C_i \in \mathcal{X}$ holds, then the raised exception matches E_i and every state path outside the exit-specific frame W_i is unchanged.
- (iv) $\mathcal{I}(\sigma)$ and $\mathcal{I}(\sigma')$ hold (invariants preserved);
- (v) $\text{derive}(\sigma.\mathcal{O}) = \sigma.\mathcal{D}$ and $\text{derive}(\sigma'.\mathcal{O}) = \sigma'.\mathcal{D}$ (derived consistency).

Definition 3 (Semantic frame). Let $\mathcal{W}$ be a set of write paths: strings of the form $\text{map}[\text{key}].\text{attr}$, $\text{map}[\text{key}]$, or channel. For a write path $p \in \mathcal{W}$ whose key is the argument attribute k , the frame checker verifies:

- $\sigma.\mathcal{O}[p] = \sigma'.\mathcal{O}[p]$ for every keyed path not in $\mathcal{W}$ (attribute stability);
- the key sets of every keyed map are equal between σ and σ' except for keyed-row write paths, which may add the args-resolved key (insertion) but never delete.

Definition 4 (Projection symmetry). A projection $\text{snap} : \text{Ctx} \rightarrow \Sigma$ is symmetric if it is applied identically before and after execution: $\sigma^{\text{pre}} = \text{snap}(c)$, $\sigma^{\text{post}} = \text{snap}(\text{exec}(c))$, and contract satisfaction is decided over $(\sigma^{\text{pre}}, a, r, \sigma^{\text{post}})$.

Definition 5 (Theory). A theory for a library L is a tuple $\langle \mathcal{A}, \mathcal{E}, S, H, \text{adapt}, \text{val} \rangle$ where $\mathcal{A}$ is the action model (covered fragment of L), $\mathcal{E}$ is the event model (trace entries), S is the pure state model, H is a stub handler interpreting actions over S and recording a trace, adapt connects H to L 's API so unmodified service code runs against the stub, and val is the validation relation (a set of proof obligations, one per program in the differential suite, each stating that the stub's observable output equals the real library's output on the covered fragment). The validation relation is established by differential testing; it is the evidence that the theory is sound, not a mere mock.

4 Contract Language

A `axiomander` contract is a first-class Python object built by the `Contract` constructor in the *fluid contract* style [?, ?]: the contract is a declarative value assembled from named clauses, not imperative assertions scattered throughout the code.

The contract language is the **pure terminating fragment of Python**: every clause is a lambda whose body may contain only arithmetic, comparison, boolean operators, attribute access, and a small fixed set of combinator calls (`extends_by_one`, `implies`, `safe_div`). Side effects, loops, exceptions, and non-boolean returns are rejected by a static purity checker before the clause is executed or lowered. Every clause must evaluate to a Python `bool`.

Definition 6 (Contract record). *A contract for an operation is $\langle \Sigma, \text{args}, \text{pre}, \text{post}, \mathcal{X}, \mathcal{I}, \mathcal{D}, \mathcal{W}, \mathcal{R}, \mathcal{G} \rangle$ where:*

- Σ is the specification state type;
- args is a typed argument record (a frozen dataclass);
- $\text{pre}(\sigma, a)$ is a boolean predicate (admissibility);
- $\text{post}(\sigma, a, r, \sigma')$ is a boolean predicate (success obligations);
- $\mathcal{X} = \{ \langle E_i, C_i, W_i, P_i \rangle \mid i \in [1..n] \}$ is the set of exceptional exits;
- $\mathcal{I}$ is the invariant set; $\mathcal{D}$ maps derived field names to pure functions; $\mathcal{W}$ and $\mathcal{R}$ are the write and read frame path sets; $\mathcal{G}$ carries ghost type, initialisation, transitions, and invariants.

4.1 Semantic frames

Frame conditions in classical DbC are verbose: In `axiomander` the contract declares only what it *may* write:

$$\mathcal{W} = \{ \text{"state.products[sku].reserved"}, \text{"state.invitations[token]}, \text{"state.reservation_log"} \}$$

The checker derives frame soundness automatically: every attribute outside $\mathcal{W}$ must be equal, every key outside the args-resolved write paths must have stable key sets, and deletions are always rejected. Keyed-row paths may *add* the args-resolved key — essential for operations such as `invite` that create rows.

4.2 Exceptional exits

Exceptions are specified as data:

$$\mathcal{X} = \{ \langle \text{InsufficientStock}, s.\text{available}(a.\text{sku}) < a.\text{quantity}, \{ \text{failure_log} \}, \text{failure_ensures} \rangle \}$$

The runner checks that when the `when` condition holds, the correct exception is raised, the exit’s frame is respected, and the failure telemetry contains exactly one new event with the declared fields. The classical “exceptions leave state untouched” discipline is the default empty frame.

4.3 Telemetry: `extends_by_one`

The combinator `extends_by_one(old, new, pred)` states that *new* is exactly *old* with one element appended, and the appended element satisfies *pred*. It is the contract language’s primitive for asserting event-channel content, used in both success and exit post-conditions. The runtime interpretation is $|new| = |old| + 1 \wedge new[-1] = old \wedge pred(new[-1])$; the lowering translates it to a list-append obligation in the generated Rocq.

5 Symmetric Projection

The same projection function bridges the concrete execution world and the abstract specification.

A Gherkin Examples row is turned into a typed witness by a pure Python function. The materializer creates a temp SQLite file from the witness, executes DDL, and inserts the initial




Figure 2: The symmetric projection. `snapshot(context)` projects the concrete world both before and after execution; contracts compare pre and post. The service is ordinary production code against SQLAlchemy; the same code runs on a real SQLite file or the theory stub (??).

rows. The runner then performs six checks per row: (1) materialize the witness, (2) pre-check invariants and admissibility, (3) execute the service and emit typed events, (4) frame-check that only declared writes changed, (5) derived-check that computed aggregates are consistent, (6) post-check the ensures clause and invariants.

The service is ordinary production code: SQLAlchemy `engine.begin()` transactions, `text()` statements, domain exceptions. Nothing in it knows about contracts, snapshots, or scenario rows.

6 Examples with Databases

We present three domains of increasing structural interest, all against SQLite through SQLAlchemy. ?? summarises their shapes and ?? gives quantitative measures.

6.1 Inventory: deltas and conditional telemetry

The inventory domain models stock reservations over a single `products` table. Contracts for `reserve`, `release`, and `restock` share the same state schema, invariants ($0 \leq \text{reserved} \leq \text{on_hand}$ for every product), and derived computations (total on-hand, reserved, available).

$$\text{pre}(s, a) = a.\text{quantity} > 0 \wedge a.\text{sku} \in s.\text{products}$$

$$\begin{aligned} \text{post}(s, a, r, s') = & s'.\text{products}[a.\text{sku}].\text{reserved} = s[\cdot].\text{reserved} + a.\text{quantity} \\ & \wedge \text{e1}(s.\text{reservation_log}, s'.\text{reservation_log}, \lambda e. e.\text{sku} = a.\text{sku}) \\ & \wedge \text{e1}(s.\text{gauge_log}, s'.\text{gauge_log}, \lambda g. g.\text{sku} = a.\text{sku} \wedge g.\text{on_hand} = s'.\text{on_hand}) \\ & \wedge \text{implies}(\text{crossed}(s, a, s'), \text{e1}(s.\text{alert_log}, s'.\text{alert_log}, \dots)) \end{aligned}$$

The alert is edge-triggered: available stock crosses from above the reorder point to at-or-below it.




Figure 3: The theory stack. The service is written once, against the real library API; the shim connects to the stub; the stub interprets actions over a pure state model and records the trace. Differential tests run the same programs against stub and real library and require identical observable behaviour.

6.2 Bank transfer: multi-key deltas and invariants

The transfer contract moves funds between two rows of one table, with a balance-nonnegativity invariant and currency-mismatch and insufficient-funds exits. Its writes span two different keyed attributes on different keys, demonstrating that frames compose over multi-delta operations.

6.3 Invitations: insertion, multi-table state, and time

The invitations domain uses insert frames, multi-table acceptance, and time-as-an-argument expiry.

$$\begin{aligned} \text{post}_{\text{invite}}(s, a, r, s') = & \quad s'.\text{invitations}[a.\text{token}].\text{status} = \text{"pending"} \\ & \wedge s'.\text{invitations}[a.\text{token}].\text{expires_at} = a.\text{now} + 7 \cdot 86400 \end{aligned}$$

with $\mathcal{W}_{\text{invite}} = \{ \text{invitations}[\text{token}], \text{invite_log} \}$ — an insert-frame operation.

$$\begin{aligned} \text{post}_{\text{accept}}(s, a, r, s') = & \quad s'.\text{invitations}[a.\text{token}].\text{status} = \text{"accepted"} \\ & \wedge s'.\text{members}[a.\text{user_id}].\text{org_id} = s.\text{invitations}[a.\text{token}].\text{org_id} \\ & \wedge s'.\text{users}[a.\text{user_id}].\text{default_org} = s.\text{invitations}[a.\text{token}].\text{org_id} \end{aligned}$$

with $\mathcal{X}_{\text{accept}}$ containing `InvitationExpired` (time check) and `EmailMismatchError` (email mismatch) exits. Token is a caller-supplied argument so the insert key is args-resolved; expiry time is an integer epoch so the comparison is pure.

7 Theories: Library Semantics as Tested Models

A *theory* for a library L is a tuple $\langle \mathcal{A}, \mathcal{E}, S, H, \text{adapt}, \text{val} \rangle$ as defined in ???. The key property is that the stub handler is validated by *differential testing*: the same program runs against the real library L and against the stub, and their observable outputs must agree on the covered fragment.

7.1 SQL theory

The SQL theory covers a deliberately small fragment: equality-WHERE `SELECT` (with ascending `ORDER BY`), single-row `INSERT`, and `UPDATE` with `SET col = col ± expr`. Statements are




Figure 4: One `SqlDomain` declaration derives runners, CLI dispatch, and the conformance suite.

parsed by `sqlglot`; anything outside the fragment raises `UnsupportedStatementError` — no silent degradation. The stub handler interprets the action model over an immutable `TableStore`, records the trace, and enforces transaction discipline. A PEP-249 DBAPI shim exposes the handler so *unmodified SQLAlchemy service code runs against the theory*. The sqlite dialect’s deferred-autobegin semantics is modelled at the connection level; `commit()` and `rollback()` are, as in real sqlite3, no-ops without an open transaction.

7.2 Logging and OpenTelemetry

The logging and OpenTelemetry theories follow the same structure. Both ship validation suites that compare stub output against the real library. The logging theory captures every `logger.info/warning/error` call as a structured `LogEmit` record; the OTel theory captures the full span lifecycle as a `SpanRecord`.

Principle 1 (Theory conformance). *A theory is admissible only if: (i) its stub and the real library agree observationally on a differential suite; (ii) its translator rejects out-of-fragment input loudly; (iii) its state model and trace are pure data. New theories follow the same checklist.*

8 Domain Declarations

Every domain in Axiomander consists of the same nine artifacts: domain types, a service, a contract, Gherkin features, witness builders, an execution context, a projection, a materializer, and a runner. Only five of these are genuinely domain-specific: types, service, contract, features, and witness builders. The remaining four — the execution context, projection, materializer, and runner — are the same infrastructure repeated verbatim per domain, with only class names and types changed.

A `SqlDomain` declaration eliminates the copied infrastructure. The domain author writes the five genuine artifacts plus a single declaration; the framework generates the rest. Concretely, the inventory domain reduces from ≈ 500 lines of boilerplate to a ~ 100 -line declaration.

The materializer creates a temporary SQLite file, executes the declared DDL, and seeds rows from the witness. The projection queries each declared table through SQLAlchemy, maps raw rows to typed domain records, and partitions the event log by channel. The runner wires these into the six-stage check pipeline. Domains with tables whose shape does not match a dataclass-to-row mapping (the key/value `limits` table in bank transfer) declare a `populate_extra` hook and query such tables manually inside the projection hook.

9 Lowering Architecture

The lowering operates in two directions. **Contract lowering** translates the Python predicate clauses into SnakeletExn specifications (FunSpecS pre/post over a two-cell heap model); this




Figure 5: The lowering pipeline: introspect, emit, score. The scoreboard has three outcomes: PROVED (machine-checked), UNKNOWN (queued for the proof oracle), and DISPROVED (a runtime counterexample is surfaced — a concrete witness of specification failure).

is implemented and all 23 obligations are machine-proved. **Implementation lowering** (§??) translates the Python service function into a SnakeletExn *program*, whose WP refinement is proved against the contract’s specification. Previously the connection between implementation and contract was established only empirically by the scenario runner (§??); with the implementation lowerer, the connection becomes machine-checked.

The contract-lowering pipeline has three stages: introspect, emit, and score.

9.1 Implementation lowering

The implementation lowerer (`impl_lower.py`) translates a Python service function body into a SnakeletExn program. It walks the Python AST and emits the corresponding SnakeletExn expression tree: assignments become Let-bindings, arithmetic and relational operators become BinOp nodes, conditionals become If nodes, and raises become Raise with structured exception payloads.

Theory-aware SQL lowering. When the parser encounters a `conn.execute(text(sql), params)` call, it delegates to the theory lowerer (`theory_lower.py`), which classifies the SQL string into a SELECT, UPDATE, or INSERT action and emits the corresponding heap-manipulating SnakeletExn expression:

- **SELECT ... WHERE key = ?** maps to a `dict_lookup_str` call with the args-resolved key;
- **UPDATE ... SET col = col + delta WHERE key = ?** produces a let-chain that loads the store dictionary, looks up the old row, computes the new field values, inserts the new row, and stores the updated dictionary back to the heap; variable-name parameters are resolved as `SVar` references, literal parameters as `SString/SInt`;
- **INSERT INTO ... (cols) VALUES (...)** emits a `dict_insert_str` call followed by a `Store`.

WP refinement proofs. The lowered program is verified against the same specification as the contract: the WP calculus proves that the SnakeletExn program refines the FunSpecS specification. All five lowering proofs are machine-checked: `AddOneLowering` (pure arithmetic, §??), `DictLowering` (heap load + dictionary lookup, 2 dict calls), `ComputeAvailableLowering` (load + 3 dict calls + subtraction), `WithdrawLowering` (load + lookup + conditional + subtraction), and `FullRestockLowering` (full restock: SELECT + null-check + conditional exception + arithmetic UPDATE + Store). All five follow the same proof pattern of `wp_bind_item` → `wp_load` → `wp_let` → `wp_dict_lookup` → `wp_binop/wp_value`.

Two-cell heap model. The lowered program operates on a two-cell heap: `store_loc` holds the products dictionary (keyed by SKU), and `trace_loc` holds the event-channel trace. A `Select` lowering loads the store dictionary and emits a `dict_lookup_str` whose key is the WHERE-column parameter; an `Update` lowering loads the old row, computes the new row via the SET expressions, and stores the updated dictionary back to `store_loc`.

```

Python source
def restock(sku, quantity):
    row = conn.execute(
        text("SELECT on_hand, reserved, reorder_point"
            " FROM products WHERE sku = ?"),
        (sku,)
    ).fetchone()
    if row is None:
        raise ProductNotFoundError(
            sku, "", quantity, "not found"
        )
    conn.execute(
        text("UPDATE products"
            " SET on_hand = on_hand + ?"
            " WHERE sku = ?"),
        (quantity, sku),
    )
    return (sku, quantity)

```

SnakeletExn lowered program

```

Let "row"
  (Let "store_d"
    (Load (Var "store_loc"))
    (Call "dict_lookup_str"
      [Var "sku"; Var "store_d"]))
  (If (BinOp EqOp (Var "row")
    (Val LitUnit))
    (Raise (Val (LitDict
      [(LitStr "_label",
        LitStr "ProductNotFoundError");
      (LitStr "0", Var "sku");
      (LitStr "1", LitStr "");
      (LitStr "2", Var "quantity");
      (LitStr "3",
        LitStr "not found")]))))
    (Let "store_d"
      (Load (Var "store_loc"))
      (Let "old_row"
        (Call "dict_lookup_str"
          [Var "sku"; Var "store_d"])
        (Let "new_row"
          (Call "row_of"
            [BinOp AddOp
              (Call "dict_lookup_str"
                [LitStr "on_hand";
                  Var "old_row"])
              (Var "quantity");
            Call "dict_lookup_str"
              [LitStr "reserved";
                Var "old_row"];
            Call "dict_lookup_str"
              [LitStr "reorder_point";
                Var "old_row"]])
          (Let "new_store"
            (Call "dict_insert_str"
              [Var "sku";
                Var "new_row";
                Var "store_d"])
            (Let "_ "
              (Store (Var "store_loc")
                (Var "new_store"))
              (Val LitUnit))))))

```

Figure 6: Implementation lowering: the `restock` service function (12 lines of Python on the left) lowers to a SnakeletExn let-chain (right) in under 10 ms. `SELECT` becomes `dict_lookup_str`, null-check becomes `BinOp EqOp/LitUnit`, structured exceptions become `LitDict` payloads, and `UPDATE` becomes a let-chain over `dict_lookup_str` → `dict_insert_str` → `Store`. Variable-name parameters (`sku`, `quantity`) are resolved as `SVar` references in the lowered form.

Fetchone and null-check. Python’s `.fetchone()` on a SQLAlchemy result is transparent in the lowered model: the `SELECT` lowering already returns the row value. A `row is None` check lowers to a `BinOp EqOp` comparison with `LitUnit`, and a subsequent `raise ProductNotFoundError` lowers to a `Raise` with a structured `LitDict` payload keyed by positional argument index.

Example: full restock service. The `restock` operation illustrates the complete lowering: `SELECT` with variable key, null check, structured exception, arithmetic `UPDATE` with store-write, and return. The Python function (left, 12 source lines) lowers to the SnakeletExn program (right) in under 10 ms.

The lowered program operates on the two-cell heap: `store_loc` for the products dictionary and `trace_loc` for the event log. In the success case the let-chain threads `store_d` → `old_row` → `new_row` → `new_store` and writes back with a single `Store`. In the missing-SKU case, the `Raise` emits a structured exception whose label and payload fields mirror the Python exception’s class name and positional arguments; a contract’s `raises` clause in the Gherkin table is lowered to a `FunSpecS` post-condition that matches the label and payload fields against the *expected* exception.

```

Values  $v \in \text{sn\_val}$ 
 $v ::= n \in \mathbb{Z} \mid b \in \mathbb{B} \mid s \in \text{string} \mid \ell \in \text{loc} \mid f \in \text{float} \mid () \mid \text{Exn}(\text{lab}, v) \mid [v_1, \dots, v_n] \mid (v_1, \dots, v_n)$ 
 $\{k_1 \mapsto v_1, \dots\} \mid \{v_1, \dots\}$ 
Expressions  $e \in \text{sn\_expr}$ 
 $e ::= \text{Val}(v) \mid x \mid \text{Let}(x, e_1, e_2) \mid \text{BinOp}(op, e_1, e_2) \mid \text{Load}(e) \mid \text{Store}(e_1, e_2) \mid \text{Alloc}(e) \mid \text{If}(e_0, e_1, e_2)$ 
 $\text{Raise}(e) \mid \text{Try}(e_1, x, e_2) \mid \text{While}(e_1, e_2) \mid \text{For}(x, e_1, e_2) \mid \text{Call}(f, [e_1, \dots, e_n])$ 
Binary operators  $op \in \text{binop}$ 
 $op ::= \text{AddOp} \mid \text{SubOp} \mid \text{MulOp} \mid \text{DivOp} \mid \text{EqOp} \mid \text{LeOp} \mid \text{LtOp} \mid \text{GtOp} \mid \text{GeOp} \mid \text{AndOp} \mid \text{OrOp} \mid \text{NeOp}$ 
 $\text{ModOp} \mid \text{InOp} \mid \text{LenOp} \mid \text{UnionOp} \mid \text{InterOp} \mid \text{AppendOp} \mid \text{LengthOp} \mid \text{DictGetOp} \mid \text{SetAddOp}$ 
 $\text{StrIndexOp} \mid \text{StartsWithOp} \mid \text{EndsWithOp} \mid \text{StrContainsOp} \mid \text{DictSetOp} \mid \text{TupleOp}$ 
Evaluation contexts  $K \in \text{list}(\text{sn\_ctx\_item})$ 
 $K_i ::= \text{LetCtx}(x, e_2) \mid \text{BinOpLCtx}(op, v_2) \mid \text{BinOpRCtx}(op, e_1) \mid \text{LoadCtx} \mid \text{StoreLCtx}(v_2) \mid \text{StoreRCtx}(e_1)$ 
 $\text{AllocCtx} \mid \text{IfCtx}(e_1, e_2) \mid \text{RaiseCtx} \mid \text{TryCtx}(x, e_2) \mid \text{ForCtx}(x, e_2)$ 

```

Figure 7: Syntax of SnakeletExn. Binary operators with invalid operand types (e.g. adding a string and an integer) are defined to raise a `TypeError` exception, matching Python.

9.2 The Rocq kernel

The Coq side comprises three layers.

SnakeletExnLang is a small imperative language in the HeapLang tradition of the Iris project, extended with first-class exceptions (following van Collem, de Vilhena and Krebbers).

The function $\text{to_val}(e) : \text{sn_expr} \rightarrow \text{option}(\text{sn_val})$ returns $\text{Some}(v)$ iff $e = \text{Val}(v)$, and None otherwise. In particular $\text{to_val}(\text{Raise}(\text{Val}(v))) = \text{None}$: a raised exception is a stuck terminal expression, not a value. The plugging function $\text{fill}(K, e)$ applies the evaluation-context stack K to e , and $\text{fill_item}(K_i, e)$ applies a single context item. **Function values** (closures, lambdas) are not yet part of the term calculus: `Call` targets statically-named functions resolved through the `FunCtx` table or through transparent Fixpoint definitions. Adding function values is future work (§??).

SnakeletExnWp provides a weakest-precondition calculus over SnakeletExn with stateful opaque specifications. The postcondition ranges over `Result := RVal v | RExn label payload`, so exception outcomes are first-class in the specification logic. A function table maps operation names to `FunSpecS` entries of the form $(pre, post)$ where $post(\sigma, vs, r, ups)$ constrains the result $r \in \{\text{RVal}(v), \text{RExn}(\ell, p)\}$ and the list of heap-cell updates ups . The kernel guarantees that each spec is total: for every pre-state σ and argument list vs satisfying $pre(\sigma, vs)$, there exist a result r and updates ups satisfying $post(\sigma, vs, r, ups)$ and respecting the map domain. All propositional proof machinery — lookup/insert lemmas (`lookup_insert_eq`, `dict_lookup_str`), update application (`apply_updates`), and invariant induction — is standard Iris reasoning over the `gen_heap` library [?, ?].

SpecPrelude provides dictionary lookup/insert helpers and their correctness lemmas, proven once rather than per contract.

9.3 Shape extraction (introspect)

The introspector (`introspect.py`) parses each clause’s Python AST into a `ContractInfo` shape object:

- **Deltas:** expressions of the form $\text{new.map}[key].field = \text{old.map}[key].field \pm \text{args.qty}$ are extracted as $(key, field, op, qty)$ tuples.
- **Scalars:** simple comparisons $(\text{args.x} > 0, \text{args.x} \neq \text{args.y})$ are reduced to (lhs, op, rhs) triples.
- **Exceptions:** the class name, conjunctive when-conditions, and payload fields are extracted.
- **Traces:** `extends_by_one` calls are detected and deferred to a later lowering pass.
- **Out-of-shape rejection:** anything not matching the supported fragment raises `UnsupportedShapeError`. Every clause is gate-checked for purity.

Pure	
$\overline{(\sigma, \text{Let}(x, v, e_2)) \rightsquigarrow (\sigma, e_2[v/x])}$	S-LET
$\overline{(\sigma, \text{If}(\text{tt}, e_1, e_2)) \rightsquigarrow (\sigma, e_1)}$	S-IF-T
$\overline{(\sigma, \text{If}(\text{ff}, e_1, e_2)) \rightsquigarrow (\sigma, e_2)}$	S-IF-F
Binary operators	
$\frac{\text{binop_eval}(op, v_1, v_2) = v}{(\sigma, \text{BinOp}(op, v_1, v_2)) \rightsquigarrow (\sigma, v)}$	S-BINOP
$\frac{\text{binop_eval}(op, v_1, v_2) \uparrow}{(\sigma, \text{BinOp}(op, v_1, v_2)) \rightsquigarrow (\sigma, \text{Raise}(\text{Exn}(\text{TypeErr}, ())))}$	S-BINOP-ERR
Heap ($\sigma(\ell)$ denotes heap lookup)	
$\frac{\sigma(\ell) = v}{(\sigma, \text{Load}(\ell)) \rightsquigarrow (\sigma, v)}$	S-LOAD
$\frac{}{(\sigma, \text{Store}(\ell, v)) \rightsquigarrow (\sigma[\ell \mapsto v], ())}$	S-STORE
$\frac{\ell' \notin \text{dom}(\sigma)}{(\sigma, \text{Alloc}(v)) \rightsquigarrow (\sigma[\ell' \mapsto v], \text{Val}(\ell'))}$	S-ALLOC
Exceptions	
$\frac{K_i \neq \text{TryCtx}}{(\sigma, \text{fill}(K_i, \text{Raise}(e))) \rightsquigarrow (\sigma, \text{Raise}(e))}$	S-RAISE
$\frac{}{(\sigma, \text{Try}(\text{Raise}(e), x, h)) \rightsquigarrow (\sigma, h[\text{Exn}(\text{lab}, p)/x])}$	S-TRY
Loops ($\text{While}(e_1, e_2) \equiv \text{For}(_, e_1, e_2)$)	
$\overline{(\sigma, \text{For}(x, \text{tt}, e_2)) \rightsquigarrow (\sigma, \text{Let}(x, \text{tt}, \text{If}(e_2, \text{For}(x, \text{tt}, e_2), ())))}$	S-FOR
$\overline{(\sigma, \text{For}(x, \text{ff}, e_2)) \rightsquigarrow (\sigma, ())}$	S-FOR-DONE
Context lifting and transparent calls	
$\frac{e = \text{fill}(K, e_0) \quad (\sigma, e_0) \rightsquigarrow (\sigma', e'_0)}{(\sigma, e) \rightarrow (\sigma', \text{fill}(K, e'_0))}$	S-LIFT
$\frac{f \in \{\text{dict.lookup_str}, \text{dict.insert_str}, \text{row.of}\} \quad \text{head_step_call}(f, [v_1, \dots, v_n], \sigma) = (\sigma', v)}{(\sigma, \text{Call}(f, [v_1, \dots, v_n])) \rightsquigarrow (\sigma', v)}$	S-CALL-T

Figure 8: Head-step relation $(\sigma, e) \rightsquigarrow (\sigma', e')$ of SnakeletExn. The lifted relation $\rightarrow$ (s-lift) plugs head steps into evaluation contexts. Transparent calls (s-call-t) are Coq Fixpoints with explicit head-step definitions and proved wp lemmas; user-level service operations are opaque, resolved through the FunCtx table of pre/post specifications.

9.4 Emission

The emitter (`emit.py`) translates a `ContractInfo` into a complete Rocq file. The generated structure per operation is:

- **Heap model:** two locations, `store_loc` (the table dict) and `trace_loc` (the event dict).
- **Row constructors:** `row.of( $f_1, \dots, f_k$ )` builds a `LitDict` with typed fields.
- **Row invariant:** `row.inv( $v$ )` holds iff there exist fields such that $v = \text{row.of}(\dots)$ and the invariant conjuncts hold (e.g. non-negativity, ordering).
- **Store invariant:** `store.inv(kvs)` lifts `row.inv` over all entries.
- **Pre/post:** `gen_pre` and `gen_post` bind all free variables existentially, look up each key's row via `dict.lookup_str`, conjunct the scalar preconditions, and (for post) describe the result and the *nested insert* of all deltas.
- **Function table:** `gen_table` maps operation names to `FunSpecS` entries, with one arm per exception exit.
- **Totality lemmas:** `gen_table_total` and `gen_table_total_pure` prove that the table satisfies the FunCtx requirements.
- **Per-delta preservation:** each delta gets a lemma showing the store invariant survives the single-row update.

9.5 Example: from Python contract to proved obligation

Consider the reserve contract from ???. The shape extractor identifies one delta (`reserved += quantity`), a scalar precondition (`quantity > 0`), one exception exit (`InsufficientStockError` with a conjunctive when-condition), and row fields of type `int`.

The Python source being lowered is the contract's post-condition:

```
lambda old_s, args, result, new_s: (
    new_s.observed.products[args.sku].reserved
    == old_s.observed.products[args.sku].reserved
    + args.quantity)
```

The mapping from Python to SnakeletExn is direct: `old_s.observed.products` becomes a lookup in the heap cell `store_loc` (a `LitDict`), `args.sku` becomes a string argument `sku`, `args.quantity` becomes an integer argument `quantity`, the arithmetic expression `s[sku].reserved + quantity` becomes a symbolic delta over row fields, and the state update becomes a `dict_insert_str` into the store dictionary. The generated Coq is:

```
Definition gen_post (sigma : sn_state) (vs : list sn_val)
  (r : Result) (ups : cell_updates) : Prop :=
exists sku order quantity store_d on_hand_sku
  reserved_sku reorder_point_sku,
  vs = [LitString sku; LitString order; LitInt quantity] /\
  sigma !! store_loc = Some (LitDict store_d) /\
  dict_lookup_str sku store_d =
    Some (row_of on_hand_sku reserved_sku reorder_point_sku) /\
  r = RVal (LitInt reserved_sku) /\
  ups = [(store_loc, LitDict (dict_insert_str sku
    (row_of on_hand_sku (reserved_sku + quantity)
      reorder_point_sku) store_d))].
```

The obligation set for this contract is:

Obligation	Statement
O1	Admissibility sanity: $\exists \sigma, vs. \text{gen_pre}(\sigma, vs) \wedge \text{store_inv}(\sigma)$ — the spec is non-vacuous.
O2	Spec consistency: $\text{gen_pre}(\sigma, vs) \implies \exists r, ups. \text{gen_post}(\sigma, vs, r, ups) \wedge \text{updates_dom_in}(\sigma, ups)$ — the success path is reachable.
O3.i	Each exception arm is satisfiable: given its pre, there exists a result and updates satisfying its post.
O4	Exit coverage: the success availability and the disjunction of all exit when-conditions partition the state space.
O5	Invariant preservation across all nested delta updates.
O8	Frame soundness: for any location $l \neq \text{store_loc}$, $\text{apply_updates}(\sigma, ups)!!l = \sigma!!l$.

The harness compiles the file against the Rocq kernel using the dependency-layer `_CoqProject` flags. Each obligation receives a three-valued verdict:

- **proved** — the obligation closes when every other lemma in the file is replaced with `Proof . Admitted.` (bisection);
- **disproved** — the obligation's negation is certified by a bundling theorem in the evidential negation layer (`_Lneg.v`), which compiles against the same kernel; a concrete runtime counterexample backs the theorem;
- **unknown** — neither proof nor disproof is certified; the obligation is queued for the proof oracle.

On a true contract this yields all PROVED; on a false contract the violated obligation is DISPROVED while sound obligations remain PROVED. The same scoreboard works for both positive and negative verdicts, enabling the dialectic loop described in §??.

9.6 Three-valued outcome and counterexample discovery

The lowering pipeline supports *three* verdicts per obligation, closing the gap between “proved” and “not proved” with a constructive counterexample model.

Evidential negation layer (`Lneg.v`). When a contract is unguarded — i.e. the implementation writes state that violates the invariant — the scenario runner discovers a runtime concrete witness: a store snapshot, argument vector, and computed bad row whose insertion breaks the invariant. The witness is optional: hand-authored (as a module-level `_WITNESS` attribute in the contract file) or implementation-yielded (through a `discover_witness()` function that runs the unguarded implementation against candidate rows and captures the first violation). The emitter materialises the witness into `Lneg.v`, an evidential negation layer containing:

- a `CounterWitness` record (store, key, arguments, bad row) as extractable data;
- computational satellite lemmas (pre-holds, update-computes, violates-inv);
- a per-obligation negation-form theorem proving that the invariant’s denial follows from a concrete row violation;
- a per-witness bundling theorem that instantiates the negation form with the candidate witness and proves it false.

The `Lneg` layer is self-contained: its `_CoqProject` maps the shared kernel, the contract definitions, and all positive layers (for the store-invariant lemma it re-proves inline). The scoreboard detects DISPROVED by checking whether the bundling theorem compiles against `coqc`.

Per-obligation compile checking via LSP diagnostics. The `rocq-piler verdict` tool uses `coq-lsp`’s whole-file diagnostic map rather than `coqc`’s first-error stop, so a failure in one obligation (e.g. `o5`) does not mask another (e.g. `o8`). The verdict is PROVED iff: (i) no severity-1 diagnostic falls inside the target lemma’s span; (ii) cross-validation via `coqc` with the target’s line range finds no error; (iii) the proof body contains no holes; (iv) `Print Assumptions` returns only the allowlisted `stdlib` and provided-context primitives; (v) the statement hash (SHA256 of the normalised statement text) matches the hash recorded in `statements.json` at emission time.

Certified disproved verdicts. The `certify_witness` tool verifies the `Lneg` layer: all satellite lemmas must compile, the negation-form theorem must close, and the bundling theorem must chain them into a witness-certified contradiction. The certified counterexample is extractable as a `CounterWitness` record pointing to the concrete store values that violate the invariant. A NOT_DISPROVED verdict means either the witness is malformed or the implementation’s violation is not covered by the current negation form.

Demonstrated end-to-end. The true reserve contract scores 6/6 PROVED across four layers. The unguarded reserve contract (missing `InsufficientStockError` guard) scores: L0 1/1 PROVED, L1 1/1 PROVED, L3 `o5` DISPROVED (certified by bundling theorem `cex_0_o5_invariant_preservation`, `o8` PROVED. The witness: `reserved = 8 → reserved = 13 > on_hand = 10`, discovered from the implementation’s own behaviour with no hand-authored witness.

The following components are **trusted** (not themselves machine-proved):

- The Rocq kernel and the `coqc` proof checker.
- The `FunSpecS` kernel and `SnakeletExnWp` calculus: hand-written Rocq that must correctly model the intended operational semantics.
- The `SpecPrelude` dictionary helpers.
- The lowering implementation itself: `introspect.py` (AST extraction and shape recognition), `emit.py` (translation to Rocq), `frames.py` (frame checking), and `purity.py` (static purity validation).

The following are **proved** by the generated obligations: admissibility non-vacuity, post-state consistency, exit coverage, invariant preservation, and frame soundness – relative to the kernel. The following are **not yet proved**: trace obligations in the lowering (enforced at runtime by the scenario runner), the correctness of the emitter itself (including the Lneg emission for counterexample witnesses), and the function-value extension of SnakeletExnLang (the current term calculus has no closures or lambdas, so first-class Python callables are not yet modelled in the lowering).

10 Evaluation

All 23 store obligations across four contracts are machine-proved. The lowering additionally emits each contract’s obligation set as dependency layers (one file per layer, plus a `_CoqProject` for parallel proof orchestration by `rocq-piler`) and a per-contract `statements.json` enforcing statement immutability. When counterexample witnesses are supplied, an evidential negation layer (`Lneg.v`) is emitted alongside the positive layers; the scoreboard classifies each obligation as `PROVED`, `DISPROVED`, or `UNKNOWN` using bisection and Lneg compilation (§??). Trace obligations (list-append post-conditions for event channels) are enforced at runtime by the scenario runner and are the active development front in the lowering.

The three example domains exercise a range of verification scenarios: single-table keyed deltas (inventory), multi-key deltas with a conservation invariant (bank transfer), and insert-frame multi-table deltas with time-based exception exits (invitations). All 42 feature rows are exercised by the generic conformance suite.

10.1 Automated Obligation Closing with `rocq-piler`

To evaluate whether the generated obligations are within reach of contemporary LLM-based proof automation, we deployed `rocq-piler` [?], an MCP server for interactive Rocq/Coq proof development, against the full obligation set of all four example contracts (release, reserve, restock, transfer).

`rocq-piler` exposes the `coq-lsp` language server through a tool interface: `focus_proof` inspects the current goal, `insert_tactics` executes a single tactic and returns the updated goal state, `check_file` verifies the complete proof against the coqc kernel, and `stratify` case-splits inductive proofs and auto-closes easy cases. Two additional tools support the three-valued outcome model: `verdict` assesses a specific obligation as `PROVED` or `NOT_PROVED` using `coq-lsp` per-range diagnostics, statement-hash immutability, hole scanning, and `Print Assumptions` allowlisting; and `certify_witness` verifies the evidential negation layer (Lneg) to certify a `DISPROVED` verdict backed by a concrete counterexample. A `_skill.md` file, emitted alongside the Coq sources (shipped with the package from a canonical template), provides Iris/gmap proof heuristics without giving specific answers.

The model (DeepSeek v4-pro) is given the full toolset, the skill guide, a 20-minute timeout per layer file, and the instruction to prove all admitted lemmas in order.

The total cost to close all 16 obligations across four contracts is approximately \$0.50–0.80 at current API pricing, with a wall-clock time of 25–35 minutes per contract in sequential mode.

Three-valued outcome demonstration. Beyond the automated proofs on true contracts, the `verdict` and `certify_witness` tools demonstrate the full dialectic cycle on an authored false contract. The unguarded reserve (no `InsufficientStockError` guard) yields: (i) L0, L1 obligations `PROVED` by `rocq-piler`; (ii) the invariant-preservation obligation (o5) is judged `NOT_PROVED` by `verdict` (LSP per-range diagnostics flag the failed proof); (iii) the evidential negation layer, containing a concrete witness ($reserved = 8 \rightarrow reserved = 13 > on_hand = 10$) discovered by running the unguarded implementation against a canonical failing row, compiles against `coqc`; (iv) `certify_witness` certifies the bundling theorem, yielding a `DISPROVED` verdict with

an extractable `CounterWitness` record. Per-obligation compile checking ensures that the false obligation (o5) does not mask the sound obligation (o8 frame soundness, which remains `PROVED`).

Several infrastructure challenges were discovered and resolved during evaluation: (i) `.vos` summary files generated by `rocq9.x` block `coq-lsp`'s import resolution and must be deleted at document-open time; (ii) the `_CoqProject` file emitted by the lowering pipeline can be corrupted by `coq-lsp` with an invalid `-R . Top` directive; (iii) `auto_admit` in `coq-lsp` corrupts proof bodies during incremental checking and must be disabled; (iv) the `set/pose` pattern for `gmap` witnesses produces proof terms that `coq-lsp` accepts but `coqc` rejects—the skill guide directs the model to use direct `exists` with map literals instead; (v) `destruct decide` on trivial equalities is opaque to the kernel and must be replaced with `rewrite decide_True`.

The hardest obligation type is the helper lemma layer (L2), which requires induction over the store list with a delta-update pattern that differs per contract (subtraction for release, addition for restock, etc.). The model succeeds when given the induction skeleton in the skill guide and enough time to iterate on the delta computation.

10.2 Implemented versus future work

The implemented system includes executable pre/post contracts, semantic frame checking, scenario verification, domain declarations with generic conformance testing, library theories for SQL, logging, and OpenTelemetry with validation suites, the contract-lowering pipeline (introspect, emit, score), and the implementation-lowering pipeline (Python AST $\rightarrow$ SnakeletExn program). All 23 store obligations are machine-proved against the Rocq kernel. All five implementation lowering proofs are verified (`AddOneLowering`, `DictLowering`, `ComputeAvailableLowering`, `WithdrawLowering`, `FullRestockLowering`).

Future work includes closing trace-cell obligations in the lowering, generalising the negation-form theorem to cover arbitrary contract shapes (currently tuned for store-invariant violations), adding function values (closures and lambdas) to `SnakeletExnLang` so that first-class Python callables can be lowered, incremental re-verification keyed on statement hashes, loop and sequence support in the implementation lowerer, and per-obligation `coqc` validation without LSP dependency.

11 Related Work

- **Design by Contract** (Meyer [?], Eiffel): provides runtime assertions without a path to mechanised proof. `axiomander` keeps the executable-contract ergonomics and adds the lowering to Rocq, plus frame inference that eliminates the verbosity of manual frame conditions.
- **Spec#** (Leino and Müller [?]): a language-integrated contract system with static verification via Boogie. Like `Spec#`, `axiomander` treats contracts as first-class; unlike `Spec#`, the contracts are not tied to the implementation language and persist independently of the code.
- **Dafny, Frama-C, KeY, Why3**: verify source-level contracts but require the source language to be the specification language. `axiomander` instead verifies an abstract model aligned with the declared frame, keeping production code unmodified and language-agnostic.
- **Iris separation logic** (Jung et al.): provides a framework for building WP calculi. `axiomander` uses the Iris `gen_heap` library as infrastructure for its Rocq kernel and shares the heap-cell update model.
- **Property-based testing** (QuickCheck [?], Hypothesis): generates random inputs and checks properties. `axiomander` contracts are simultaneously the property oracle and the proof input; the Gherkin tables provide the concrete test data.

- **Behaviour-driven development** (Gherkin [?]): specifies behaviour in scenario tables. `axiomander` connects Gherkin to executable contracts and formal proof, so the scenarios become verified test data.
- **TLA⁺**: specifies systems in temporal logic. `axiomander` targets operational semantics and explicit state transitions with a considerably lower barrier to entry (Python lambdas vs. TLA⁺ formulas).

Internal validity. The lowering pipeline is trusted code: the emitter’s translation from `ContractInfo` to Rocq is not itself proved correct. The frame checker and purity validator are similarly trusted. The semantic object of the proof is the contract, not the implementation — the scenario runner tests the latter empirically.

External validity. The three example domains are small but representative of database-backed services: they exercise single-table, multi-table, insert-bearing, and time-bearing patterns. The theory mechanism has been demonstrated for three libraries (SQL, logging, OTel) but the approach to adding new theories is manual.

Construct validity. The claim that contracts form a “semantic intermediate language” is currently instantiated via a specific AST extraction and lowering pipeline. The contract fragment is restricted to what the introspector and emitter can handle; expanding the fragment is future work.

12 Conclusion

This paper has presented the thesis that executable behavioural contracts form a semantic intermediate language from which testing, runtime checking, differential validation, and machine-checked proof obligations can all be derived. A vericoding workflow leaves the specification as the lasting artifact and the code as a replaceable implementation detail.

The `axiomander` toolchain operationalises this thesis through seven contributions: contracts as semantic intermediate representation, semantic frame inference, differentially validated library theories, a symmetric projection architecture, automatic lowering to Rocq proof obligations, declarative domain specification, and a three-valued outcome model with constructive counterexample discovery. The evaluation is empirical and machine-checked: 23 store obligations across four contracts are proved against a hand-verified WP calculus, and the implementation lowerer produces `SnakeletExn` programs whose WP refinement is fully proved (5 of 5 complete). Theory-aware SQL call lowering bridges the gap between Python `SQLAlchemy` code and heap-manipulating `SnakeletExn` expressions, enabling end-to-end verification from Python source to Rocq proof.

References

- [1] Bertrand Meyer. *Applying “Design by Contract”*. IEEE Computer, 25(10):40–51, 1992.
- [2] K. Rustan M. Leino and Peter Müller. *Using the Spec# Language, Methodology, and Tools to Write Bug-Free Programs*. In LASER Summer School 2008/2009, LNCS 6029, Springer, 2010.
- [3] Gavin Mendel-Gleason. *axiomander: Specification-Driven Verification*. Source code and documentation. <https://github.com/scidonia/axiomander>, 2025–2026.
- [4] Cucumber Ltd. *Gherkin Reference*. <https://cucumber.io/docs/gherkin/reference/>, 2025.

- [5] Koen Claessen and John Hughes. *QuickCheck: A Lightweight Tool for Random Testing of Haskell Programs*. In ICFP 2000.
- [6] John C. Reynolds. *Separation Logic: A Logic for Shared Mutable Data Structures*. In LICS 2002.
- [7] Scidonia. *SnakeletExn and FunSpecS: A Small Imperative Language and Weakest-Precondition Specification Kernel*. Coq development, 2026.
- [8] Ralf Jung, Robbert Krebbers, Jacques-Henri Jourdan, et al. *Iris: Monoids and Invariants as an Orthogonal Basis for Concurrent Reasoning*. In POPL 2018.
- [9] Ralf Jung. *The Iris Project*. <https://iris-project.org/>, 2018.
- [10] Tom Van Collem, Bas de Villhena, and Robbert Krebbers. *Exceptions in the Snakelet Language*. Coq development, 2026.
- [11] The Iris Project. *HeapLang: A Simple Imperative Language for Verification*. <https://iris-project.org/>, 2026.
- [12] Jeremy Tobin-Hochstadt and Matthias Felleisen. *The Design and Implementation of Typed Scheme*. In POPL 2008.
- [13] Patrick Cousot and Radhia Cousot. *Abstract Interpretation: A Unified Lattice Model for Static Analysis of Programs by Construction or Approximation of Fixpoints*. In POPL 1977.
- [14] William McKeeman. *Differential Testing for Software*. Digital Technical Journal, 10(1):100–107, 1998.
- [15] Jean-Christophe Filliâtre. *Why3: A Platform for Deductive Program Verification*. In CAV 2013.
- [16] Xavier Leroy et al. *CompCert: Formally Verified Optimizing Compiler*. <https://compcert.org/>, 2016.
- [17] Zoe Paraskevopoulou et al. *QuickChick: Property-Based Testing in Coq*. JFP, 2022.
- [18] Grigore Roşu. *From Rewriting Logic, to Programming Language Semantics, to Program Verification*. In LOPSTR 2013.
- [19] Scidonia. *rocq-piler: An MCP Server for Interactive Rocq/Coq Proof Development*. <https://github.com/scidonia/rocq-piler>, 2026.

Stage	What happens
1. Feature	A Gherkin <code>reserve.feature</code> contains: Given a product "S1" with on-hand 100 reserved 10 ... with an Examples table of concrete rows.
2. Witness	A pure Python function maps each row to a typed <code>ReserveArgs(sku="S1", order_id="01", quantity=30)</code> and an initial products dict: <code>{"S1": Product(on_hand=100, reserved=10, ...)}</code> .
3. Materialize	A temp SQLite file is created, DDL run, and rows inserted. The concrete world is: <code>ExecutionContext(engine=create_engine("sqlite:///tmp/..."), events=EventLog())</code> .
4. Pre-snapshot	<code>snapshot(context) → SpecState</code> : reads products via SQLAlchemy, partitions the empty event log, computes derived totals, and returns an immutable frozen dataclass.
5. Admissibility	<code>requires(state, args)</code> is evaluated: <code>args.quantity > 0 ∧ args.sku ∈ state.observed.products</code> . If <code>outcome = "rejected"</code> , admissibility is expected to fail; for "success", it must hold.
6. Invoke service	<code>InventoryService().reserve(engine, "S1", "01", 30)</code> runs against the real SQLite file through SQLAlchemy. On success it returns a <code>ReservationReceipt</code> ; the effect-emitting wrapper records <code>StockReserved</code> and gauge events in the context's event log.
7. Post-snapshot	The same <code>snapshot</code> function runs against the now-mutated SQLite file, producing the post-state <code>SpecState</code> with updated reserved count and appended event tuples.
8. Contract check	The runner evaluates: (a) frame — everything outside <code>writes</code> unchanged; (b) derives — recomputed aggregates consistent; (c) ensures — reserved increased by exactly the quantity, exactly one <code>StockReserved</code> event emitted with exact fields; (d) invariants — <code>reserved ≤ on-hand</code> , both non-negative.
9. Lower	The same <code>Contract</code> object is introspected by <code>axiomander.lower</code> : deltas, scalars, and exception arms are extracted; a Rocq obligation file is emitted. <code>coqc</code> compiles it; per-obligation bisection yields a scoreboard.
10. Prove	<code>coqc -R coq " coq/gen/GenReserveObligations.v</code> succeeds. O1–O8 are machine-checked: admissibility sanity, spec consistency, exit coverage, invariant preservation, and frame soundness — all PROVED.

Table 1: End-to-end walkthrough of the reserve contract. Ten stages, one spec, three levels of assurance.

Domain	Tables	Operations	Feature rows	Obligations
inventory	1	reserve, release, restock	22	6/6 · 6/6 · 4/4
bank_transfer	2	transfer	11	7/7
invitations	3	invite, accept	9	(runtime verification)

Table 2: The three database-backed example domains. All feature rows are exercised by the conformance suite; 23 store obligations across 4 contracts are machine-proved.

Metric	Value
Python source lines (axiomander core + examples)	$\approx 11,000$
Rocq source lines (kernel + proofs)	$\approx 5,000$
Number of example domains	3
Number of contracts	6
Number of feature rows exercised	42
Generated obligations per contract	4–8
Store obligations proved	23/23 (automated over 4 contracts)
Trace obligations (runtime only)	deferred
Implementation lowering proofs	5/5 (AddOne, Dict, ComputeAvailable, Withdraw, FullRestock)
Counterexample WPs (Lneg)	1/1 (reserve)
Coq interpreter + partial evaluator	compiles (5 sanity checks pass)
BDD lowering tests (SnInterp)	13 (9 lowering + 4 counterexample)
Runtime verification per row	< 0.1 s
Lowering time (introspect + emit)	< 0.5 s per contract
Python-to-SnakeletExn lowering time	< 10 ms per function
Rocq compilation per obligation file	< 30 s

Table 3: Quantitative measures of the `axiomander` implementation. All store obligations are machine-proved.

Contract	L0	L1	L2	L3
release	✓	✓	✓	✓
restock	✓	✓	✓	✓
reserve	✓	✓	✓	✓
transfer	✓	✓	✓	✓

Table 4: Automated proof results: all 16 obligations across four contracts are closed by `rocq-piler` with `DeepSeek v4-pro`. L0 = admissibility, L1 = spec consistency, L2 = helper lemmas, L3 = preservation + frame soundness.

Obligation Type	Solved	Avg Time	Avg Cost
Admissibility (L0)	4/4	591 s	\$0.036
Spec Consistency (L1)	4/4	307 s	\$0.020
Helper Lemmas (L2)	4/4	536 s	\$0.026
Preservation (L3)	4/4	430 s	\$0.030

Table 5: Cost and timing per obligation type. Spec consistency is the cheapest (one-line `gen_table_total` lemma); admissibility the slowest (witness construction with `gmap singletons`).